DEEP DIVE INTO SEARCH SYSTEMS

HI!

- Applied Research Engineer - Spam and Relevance team, LinkedIn
- Software Engineer - Search team, LinkedIn
- Undergraduate Thesis - Search team, LinkedIn
- Systems and Network Engineering Intern, LinkedIn
- BITS Pilani - B.E. *(Hons)* Computer Science

GLOBAL INTERNET TRAFFIC IN 2013 WAS APPROXIMATELY 5,000,000,000,000,000 BYTES (5 BILLION GIGABYTES).

CHARACTERISTICS (V'S) OF BIG DATA

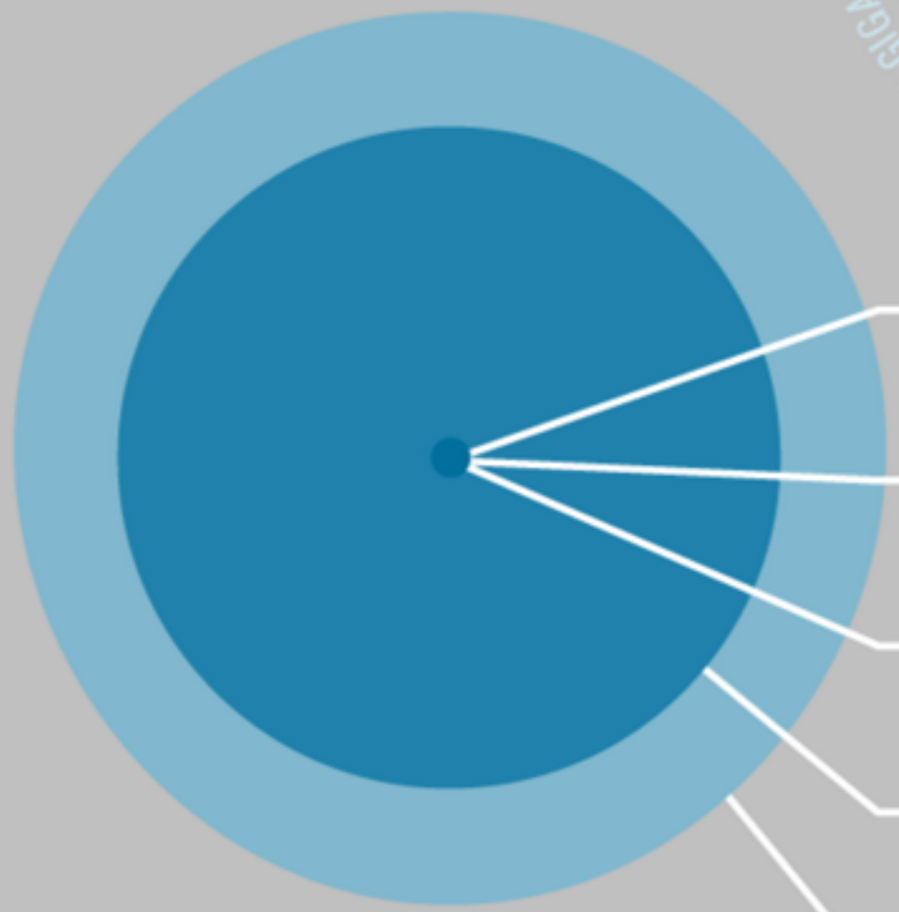

1992
100GB/DAY

1997
100GB/HOUR

2002
100GB/SECOND

2013
28,875GB/SECOND

2018
50,000GB/SECOND

Global internet population
GREW 14.3% BETWEEN
2011 & 2013

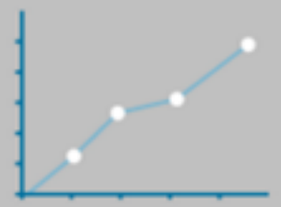

3 BILLION
The number of people who have access to the internet today equals that of the world's population in 1960

Source:
vcloudnews

DATA EXPLOSION

- 90% of the world data is created in last two years alone
- The Indexed Web contains at least 4.72 billion pages
- Why I am talking about this?
  - Without retrieval this data has very little use
- Imagine a world without Search Engines

INFORMATION RETRIEVAL

- Lets define the problem

  Given a query, find material (documents) of an unstructured nature (usually text) from a large collection (corpus)

- Not just web search!
  - E-mail search
  - Search on your laptop
  - Search on Amazon, LinkedIn

- First step?
  - Collecting Documents!

- What happens after all the documents are collected?
- They are stored in the form of a “Search Index”

SEARCH INDEX

- How should the documents be stored to ensure fast retrieval?
  - Just store all the documents in sequence and go through them
    - Called *grep*
    - More than 4.7 Billion documents on web!
    - 450 M members on LinkedIn
      - 450 M x 1 ms = 125 hours!
- Any guesses?
  - Hint - How do you search for a term in a book?

SEARCHING FOR TRAVELING SALESMAN?

- Inverted Index!
  - Terms are matched against documents and not vice-versa

INVERTED INDEX

ID	Text
1	Summer is hot. Mangoes are ripe during summer.
2	Summer is warm here
3	Why is summer hot here
4	Later we found why

Term	Document IDS
Mangoes	1
are	1
ripe	1
during	1
summer	1, 2, 3
is	1,2, 3
warm	2
here	2, 3
why	3, 4
hot	1, 3
later	4
we	4
found	4

QUERY

- A query is interpreted as a Boolean Query
  - Examples on LinkedIn Recruiter search -
    - html AND css
    - java OR python
    - implicit - synonyms
- manager NOT software

- Boolean Retrieval
  - OR / AND
    - term1 - 1,2,3,5,6 | term2 - 1,3,4,5,7,8
    - **MERGING**
      - term1 OR term2 - 1,2,3,4,5,6,7,8
      - term1 AND term2 - 1,3,5
- Time complexity
  - O[m+n]

Dictionary

Posting List

Term	Freq	Document IDS
Mangoes	1	1
are	1	1
ripe	1	1
during	1	1
summer	3	1, 2, 3
is	3	1,2, 3
warm	2	2
here	2	2, 3
why	2	3, 4
hot	2	1, 3
later	1	4
we	1	4
found	1	4

- What Data Structure is used?
  - Posting list -
    - Variable length array vs Linked List
      - Frequency of updates
      - Space and time overhead
- Accessing terms -
  - Sorted array
  - Hash table
  - B+ trees

THREE PILLARS OF SEARCH

- Any search infrastructure will expose APIs for these three functionalities -
  - Building Index
    - Preprocessing and tokenization
  - Retrieval
  - Ranking

PREPROCESSING

- Document Preprocessing (part of index build)
  - Tokenization for index creation
    - Apply filters like - lowercase, removing punctuations, stopwords
    - Stemming
  - Calculating TF and IDF values for TF-IDF scoring
- Query Preprocessing (part of retrieval)
  - Query tagging
  - Tokenization (should be similar to document preprocessing)
  - Query expansion
    - Synonym expansion

PREPROCESSING

- Stemming

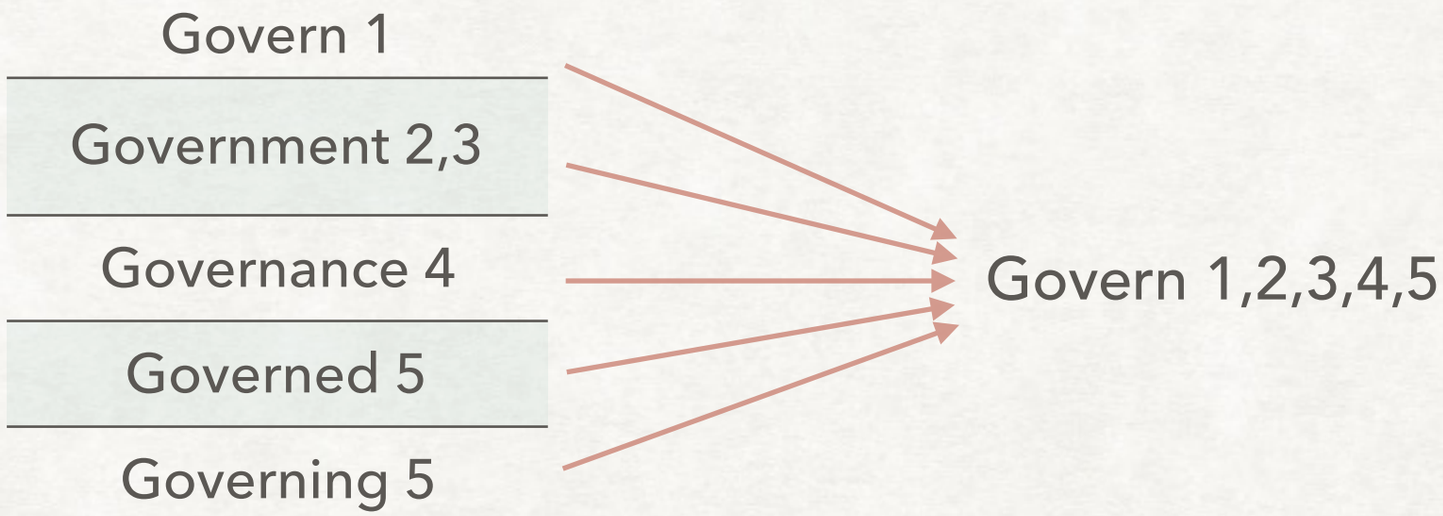

- Synonym Expansion

"slow router problem"

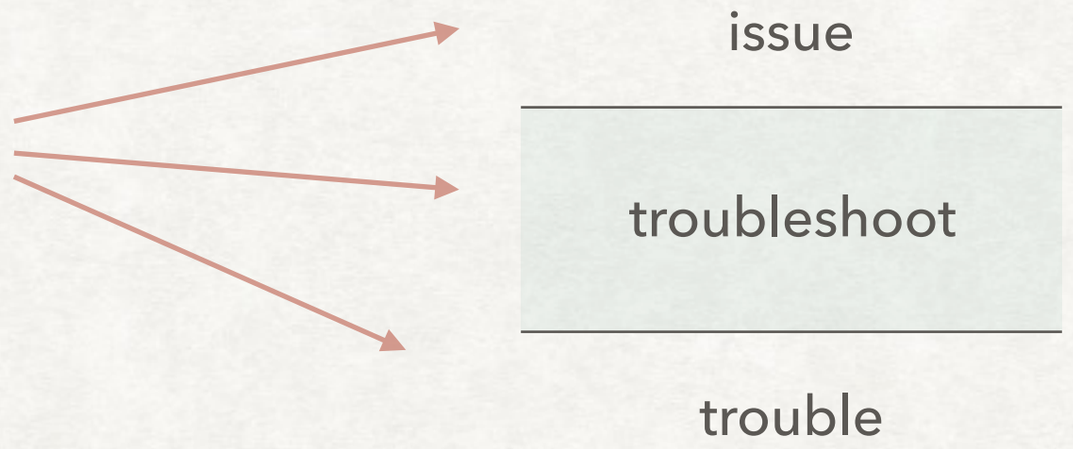

RETRIEVAL

Term	Document IDS
Mangoes	1
are	1
ripe	1
during	1
summer	1, 2, 3
is	1,2, 3
warm	2
here	2, 3
why	3, 4
hot	1, 3
later	4
we	4
found	4

Query - hot summer

Sample rewritten queries -

- hot AND summer
- (hot OR warm) AND summer
- hot OR summer

STATIC RANK & EARLY TERMINATION

- A global score of the document
- Each document has one SR but multiple documents can have same SR
- Could be anything from number of connections / followers; length of the documents; Social signals; Page Rank
- Used in early termination
  - Number of documents to score
- Posting Lists are sorted on the basis of Static Rank

RANKING / SCORING

- Retrieved documents are scored based on the combination of -
  - Query dependent features
    - Term Frequency,
    - Min window match,
    - Exact match in title, etc.
  - Query independent features
    - Length of the document,
    - Social Signals,
    - Number of followers of the author, etc.

- Once we have the ranked list of document IDs, we obviously won't show them to the users!

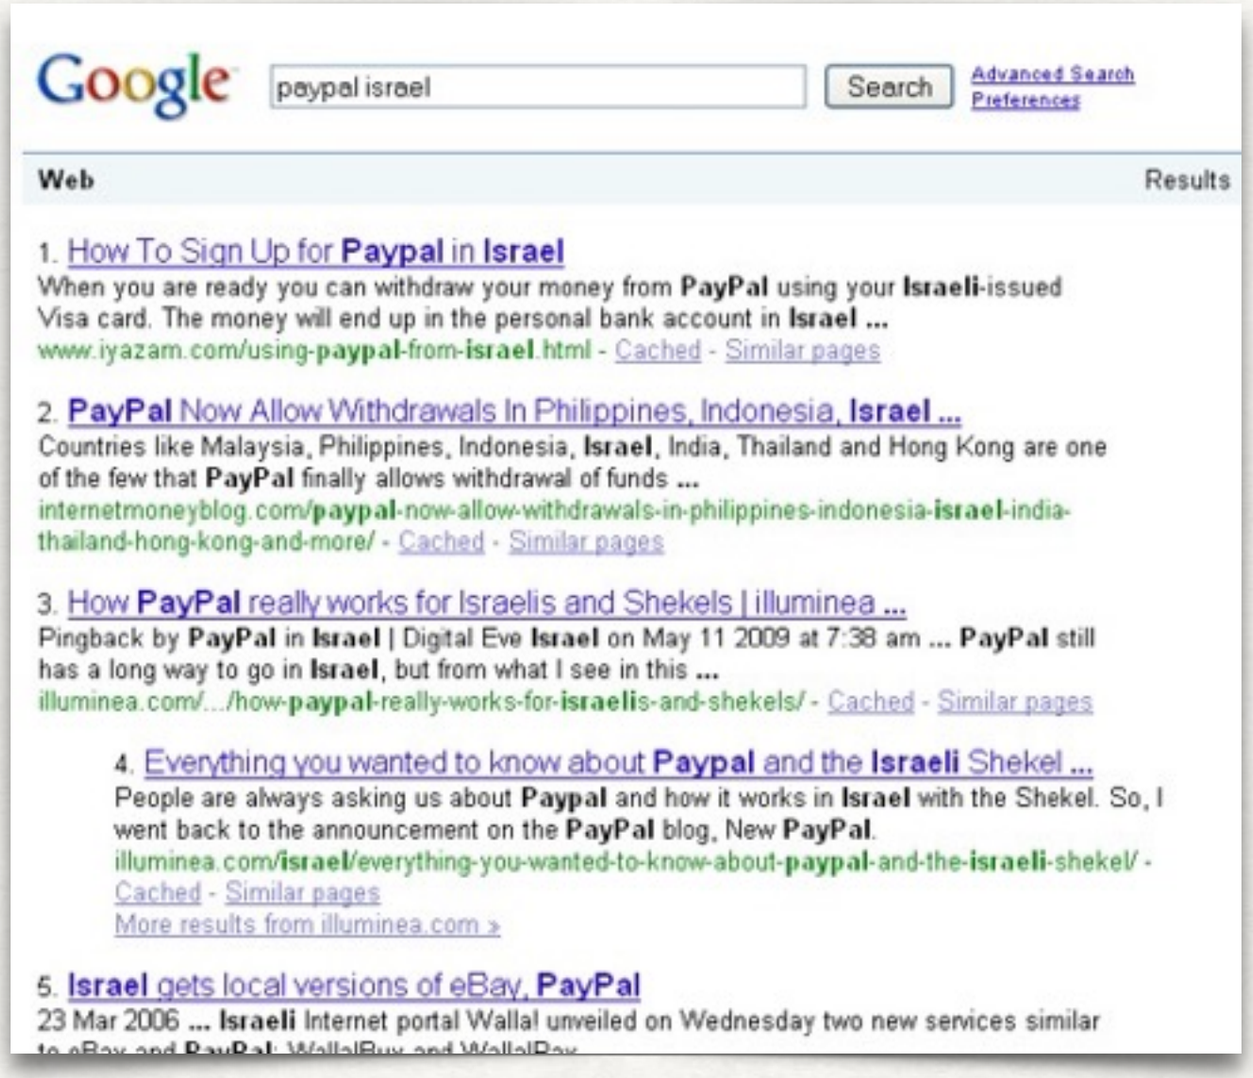

- Title, snippet, URL
- How are they stored?

	title	snippet
1	Summer	Hot summer
2	warm	warm in summer
3	hot	summer
4	later	later why

Term	Document IDS
Mangoes	1
are	1
ripe	1
during	1
summer	1, 2, 3
is	1,2, 3
warm	2
here	2, 3
why	3, 4
hot	1, 3
later	4
we	4
found	4

FORWARD INDEX

- Mapping of document ID to forward fields (title / snippet)

	title	snippet
1	Summer	Hot summer
2	warm	warm in summer
3	hot	summer
4	later	later why

Forward Index!

Term	Document IDS
Mangoes	1
are	1
ripe	1
during	1
summer	1, 2, 3
hot	1
is	1,2, 3
warm	2
here	2, 3
why	3, 4
hot	1, 3
later	4
we	4
found	4

BASIC SEARCH CONCEPTS

- Document
- Query
- Boolean Retrieval model
- Inverted Index
- Forward Index
- Retrieval
- Ranking

SCALING

- All the data won't fit in one computer
- Even if it does, we should leverage parallelization
- How to break the index into multiple machines?
  - Shards!
- Let's consider inverted index first

Simple Master Slave Configuration

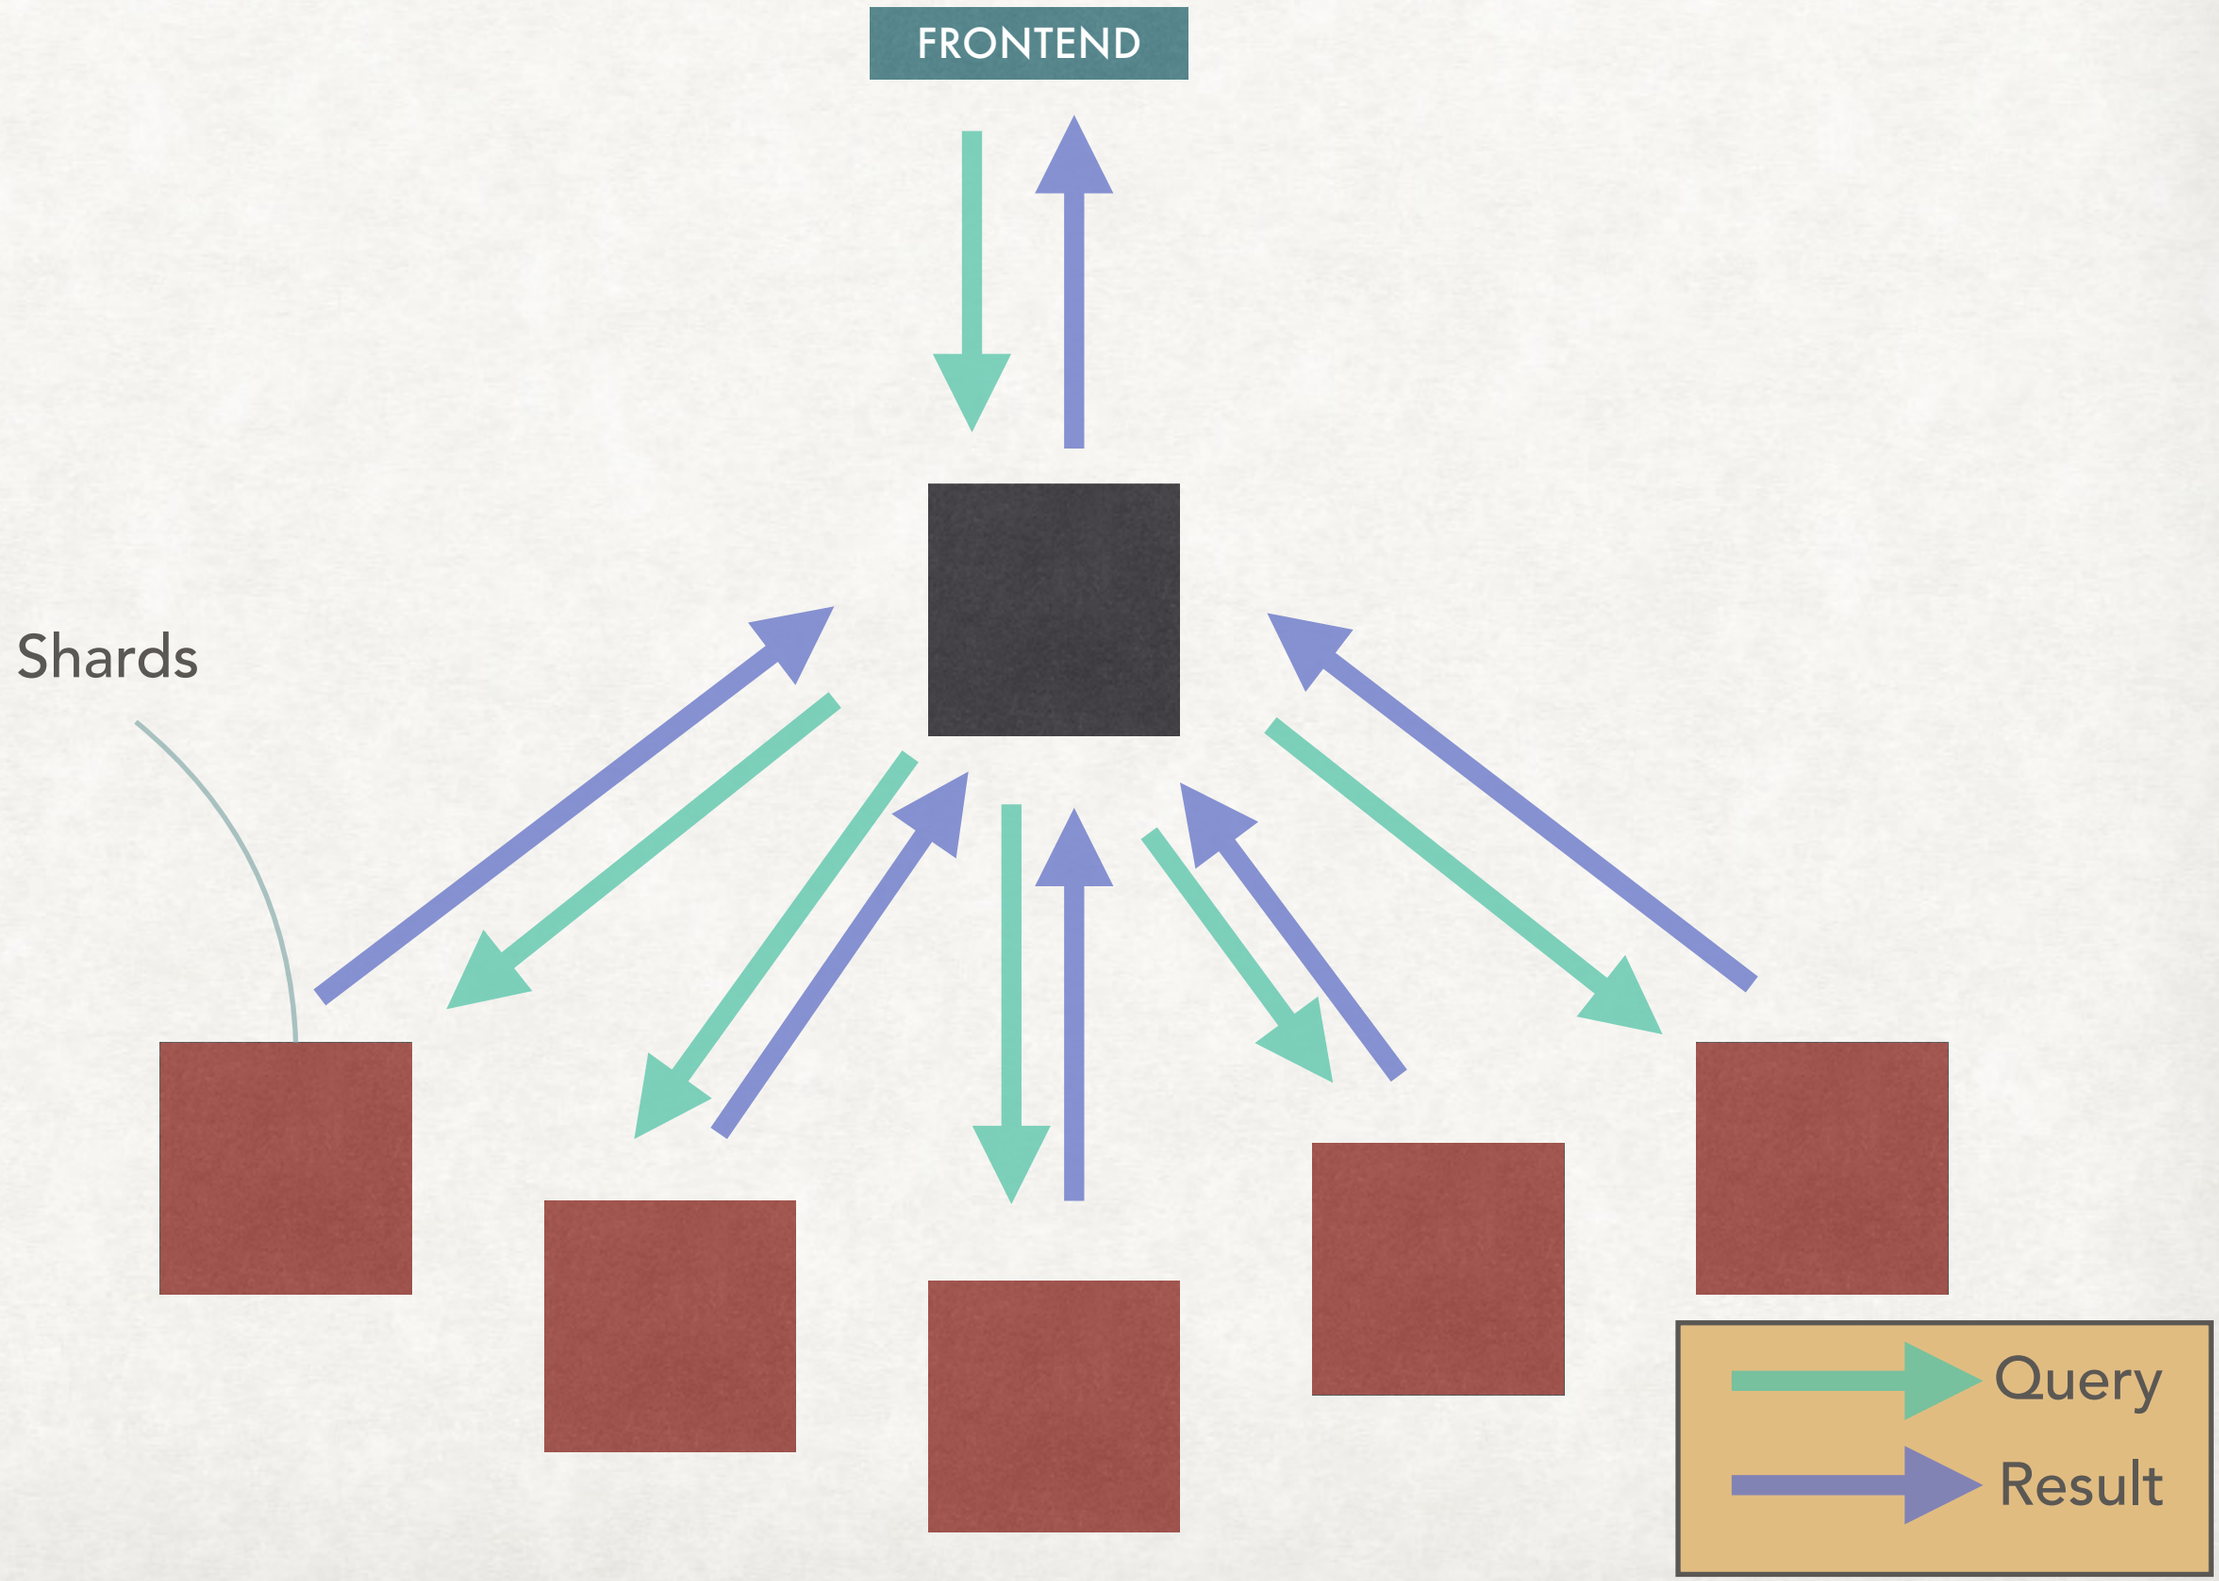

SCALING

- Document vs Term based partitioning?
  - Failover of a machine
  - Latency
  - Leveraging Parallelization
  - Maintaining forward index
  - All the above favor document based partitioning!

Simple Master Slave Configuration

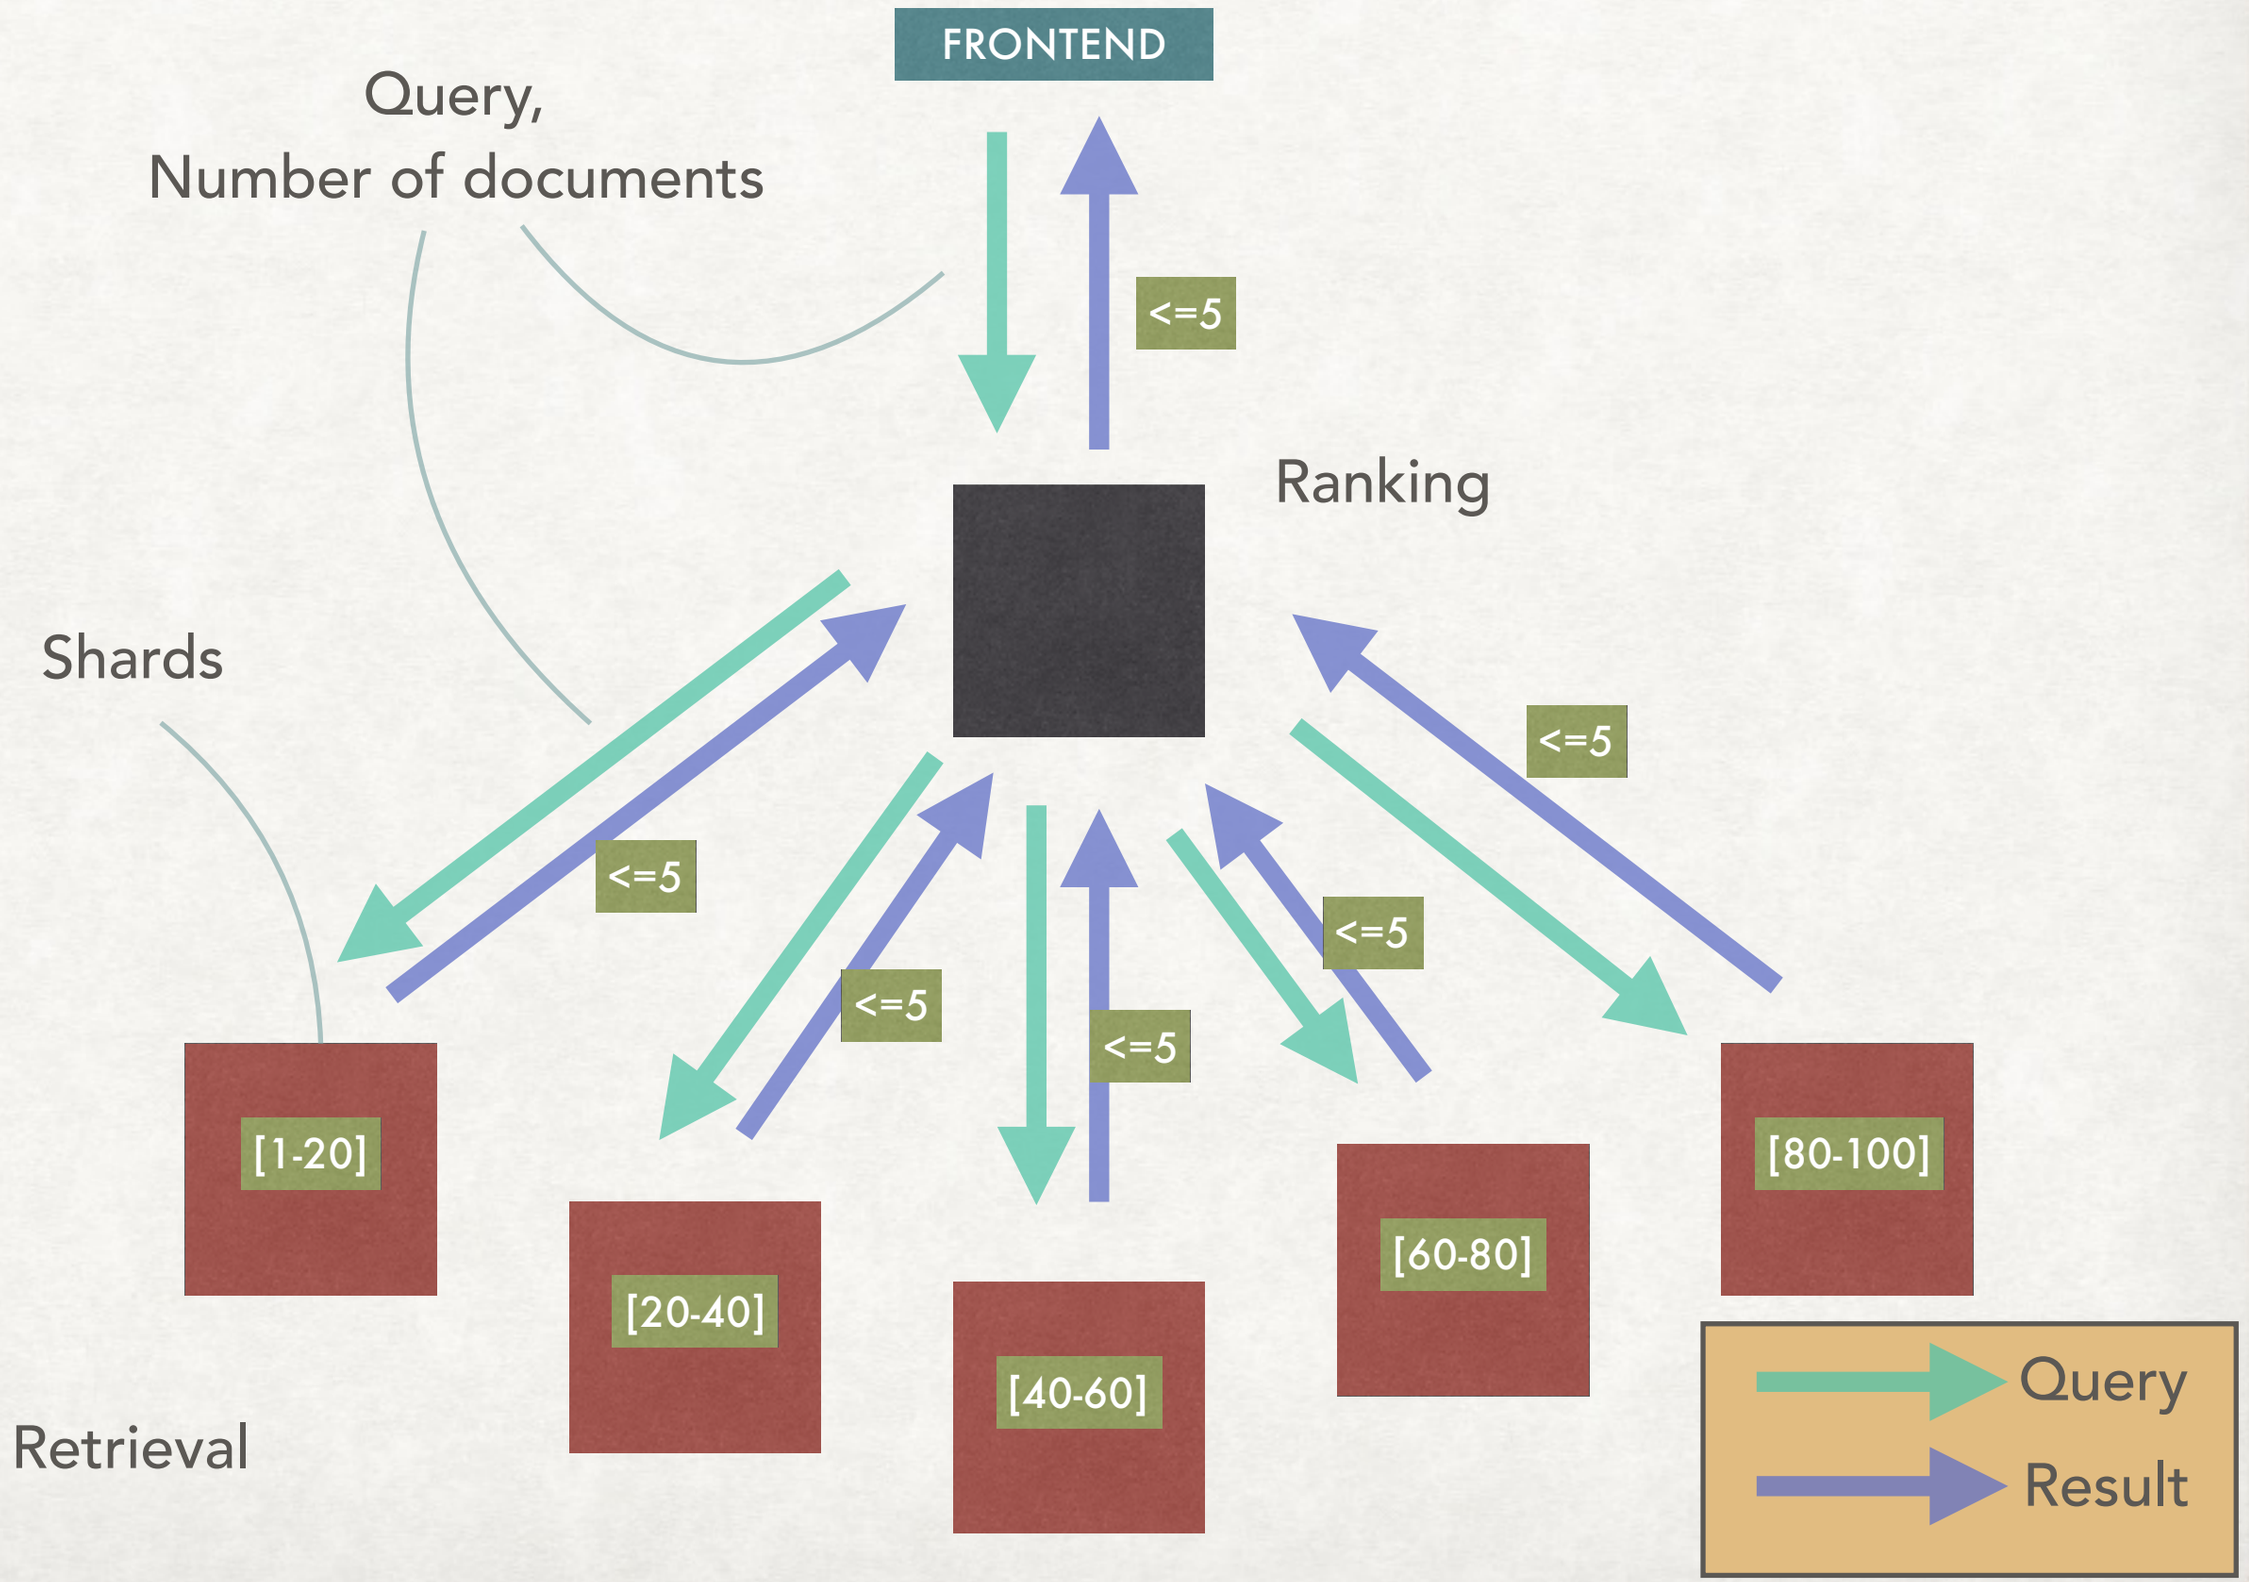

Query - hot OR summer
Number of documents - 2

Term	Document IDS
Mangoes	1
are	1
ripe	1
during	1
summer	**1, 2**
hot	1
is	1,2
warm	2
here	2
hot	**1**

1, 2

Term	Document IDS
summer	3
is	3
here	3
why	3, 4
hot	3
later	4
we	4
found	4

3

LINKEDIN SEARCH STACK (GALENE)

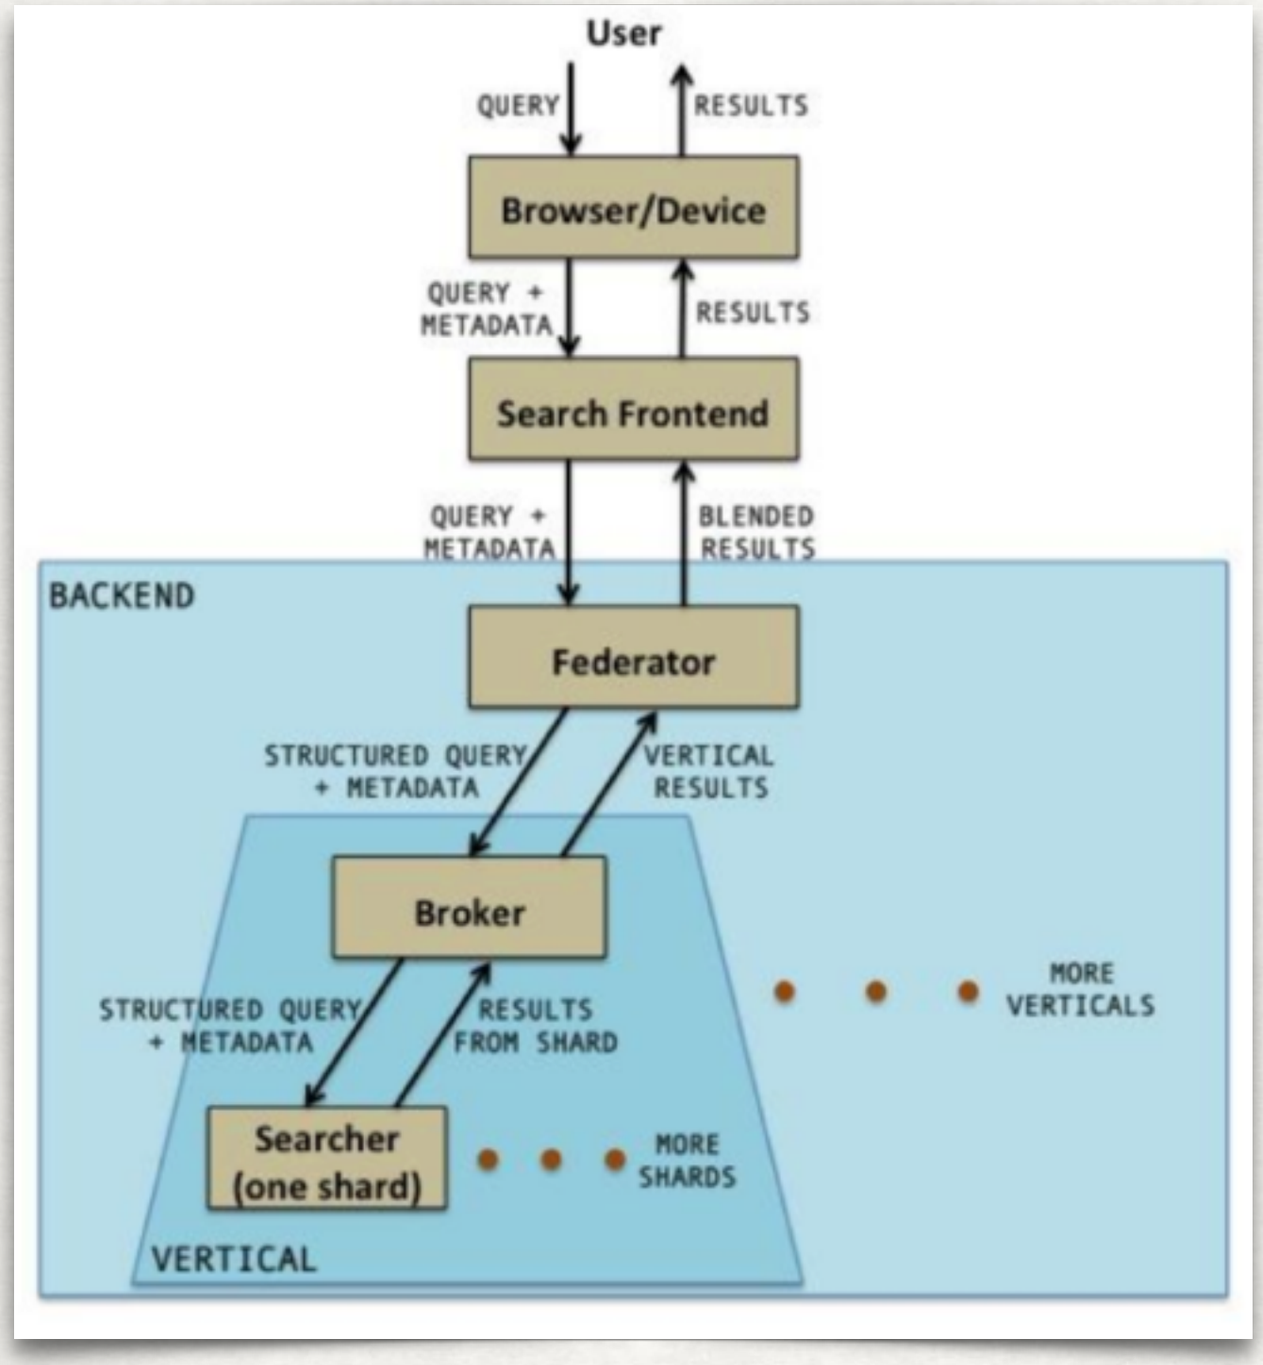

- Federator and Broker
  - Rewrites the Query
  - Fans out
  - Merges the results
  - Requires num of docs to return

- Searcher
  - Operates on single shard
  - Takes rewritten query and retrieves the documents
  - Scores the documents - using query, input metadata, match info
  - Requires num of docs to score

QUESTIONS?

Mansi Gupta

mgupta1410@gmail.com